

Stirring Up Neural ODEs: A Chemical Engineer's perspective

Prithvi Dake

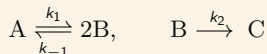
Department of Chemical Engineering




ME255NN Report
March 20, 2025

Gentle intro with a simple problem

Consider following *true* model of linear chemical kinetics taking place in a well-mixed batch reactor with $c = (c_A, c_B, c_C)^T$ and initial concentration c_0 :



$$\begin{bmatrix} \dot{c}_A \\ \dot{c}_B \\ \dot{c}_C \end{bmatrix} = \begin{bmatrix} -k_1 & k_{-1} & 0 \\ 2k_1 & 2k_{-1} - k_2 & 0 \\ 0 & k_2 & 0 \end{bmatrix} \begin{bmatrix} c_A \\ c_B \\ c_C \end{bmatrix}$$




Training a Physics-Informed Neural Network (PINN) ¹

- 1: **Input:** Neural network architecture f_{NN}
- 2: **Physical Model:** $Lx = f$ $B_1x = 0$ $B_2x = 0$ ▷ True or known physics
- 3: **Initialize** neural network parameters θ
- 4: **repeat**
- 5: $x = f_{NN}(x, t; \theta)$
- 6: $V_{ODE} = |Lx - f|^2$ ▷ Prediction error minimization
- 7: $V_{BC} = |B_1x|^2 + |B_2x|^2$ ▷ Enforce boundary conditions
- 8: $V = \lambda_1 V_{PDE} + \lambda_2 V_{BC}$ ▷ Combine loss terms
- 9: $\theta \leftarrow \theta - \eta \nabla_{\theta} V$ ▷ Update θ using gradient descent
- 10: **until** convergence criteria is met
- 11: **Return** trained model f_{NN}

¹Raissi et al. (2019)

Background: If PINNs why NODEs?



- I view PINNs as a softer version² of *method of weighted residuals*.³
- Mostly used as surrogate numerical solver for PDEs.
- Can learn some physics but has no *structure*.
- Causality problem (no sequential model, future depends on past).⁴

Recent advances

Physics-informed neural ODE (PINODE)⁵

Lift-and-embed: Koopman operators⁶

²By softer I mean PINNs minimize rather than exactly satisfy equality constraints

³Villadsen and Stewart (1967)

⁴Wang et al. (2022)

⁵Sholokhov et al. (2023)

⁶Liu et al. (2024)

A state-space representation: Euler discretized system

$$x^+ = x + f_{NN}(x, u; \theta)$$


$$y = g_{NN}(x, u; \theta)$$

$$x^+ = f_{NN}(x, u; \theta)$$

$$y = g_{NN}(x, u; \theta)$$

Parallels with FIR and ARX models in system identification ⁷.

Backprop strategy: Brute-force chain rule (store intermediate states).



$$\text{If } x \in \mathbb{R}^n, \theta \in \mathbb{R}^p, \\ S = \frac{\partial x}{\partial \theta^T} \in \mathbb{R}^{n \times p}$$

⁷Ljung (1999)

Don't want to store intermediate states? Forward sensitivities for IVPs ⁹

- Consider a non-linear state-space model ($y = x$):

$$\dot{x} = f(x, u; \theta) \quad x(0) = g(x_0; \theta)$$

- Sensitivity is defined as and evolves as ⁸:

$$S = \frac{\partial x}{\partial \theta^T}$$
$$\frac{dS}{dt} = \frac{d}{dt} \left(\frac{\partial x}{\partial \theta^T} \right) = \frac{\partial}{\partial \theta^T} \left(\frac{dx}{dt} \right) = \frac{\partial}{\partial \theta^T} f$$

- Augmented system of ODEs:

$$\begin{bmatrix} \dot{x} \\ \dot{S} \end{bmatrix} = \begin{bmatrix} f \\ \frac{\partial f}{\partial x^T} S + \frac{\partial f}{\partial \theta^T} \end{bmatrix} \quad \begin{bmatrix} x(0) \\ S(0) \end{bmatrix} = \begin{bmatrix} x_0 \\ \frac{\partial g}{\partial \theta^T} \end{bmatrix}$$

- Notice again: $S \in \mathbb{R}^{n \times p}$ i.e sensitivities we need to solve for increase **linearly** with p .
Imagine how bad this could get for FNN with millions of parameters and states!

⁸Using Clairaut's Theorem on Mixed Partialials

⁹Derivation in Rawlings and Ekerdt (2020)

How to avoid the linear cost? Pontryagin's adjoint method



Lev Semenovich Pontryagin

P. S. Alexandrov

Andrey N. Kolmogorov

- Consider following optimization problem and the Lagrange functional given data \tilde{x} :

$$\begin{aligned}\min_{\theta} V(x) &= \int_0^t |x - \tilde{x}|^2 dt = \int_0^t g(x) dt \\ \min_{\theta, \lambda} \mathcal{L} &= \int_0^t g dt + \int_0^t \lambda(t)^T \left(f - \frac{dx}{dt} \right) dt\end{aligned}$$

- The augmented system of ODEs is ¹⁰:


$$\begin{bmatrix} \dot{x} \\ \dot{\lambda}^T \end{bmatrix} = \begin{bmatrix} f_{NN} \\ -\lambda^T \frac{\partial f_{NN}}{\partial x^T} \end{bmatrix} \quad \begin{bmatrix} x(t) \\ \lambda^T(t) \end{bmatrix} = \begin{bmatrix} x_t \\ -\frac{\partial g}{\partial x^T} \Big|_t \end{bmatrix}$$

- The gradient of loss wrt parameters is calculated as:

$$\frac{\partial V}{\partial \theta^T} = \lambda^T(0) \frac{\partial x}{\partial \theta^T} \Big|_0 + \int_0^t \lambda^T \frac{\partial f}{\partial \theta^T} dt$$

¹⁰Refer Sengupta et al. (2014) for paper and <https://youtu.be/k6s2G5MZv-I?si=ZIyP7MvgzTK7D9kn> for a walkthrough. Additionally, I have worked out the derivation step-by-step in my report.

An ultimate panacea?




- Consider $\dot{\lambda} = \pm 10\lambda$ and $\lambda(0) = 0$
- We see that the adjoint variable will either *explode* or *vanish*. This heavily impacts the learning problem.
- The Jacobian $\frac{\partial f}{\partial \mathbf{x}^T}$ can be ill-conditioned.
- We solve the adjoint equation by quadrature, thus numerical noise can still be a problem.

Remedies

- Don't solve adjoint over long horizons, but rather small windows and store the initial conditions for those intermediate windows.
- Popularly called as *checkpointing*¹¹.

¹¹Zhuang et al. (2020)

- 1: **Input:** Neural network architecture f_{NN}
- 2: **ODE:** $\begin{bmatrix} \dot{x} \\ \dot{\lambda}^T \end{bmatrix} = \begin{bmatrix} f_{NN} \\ -\lambda^T \frac{\partial f_{NN}}{\partial x^T} \end{bmatrix} \quad \begin{bmatrix} x(t) \\ \lambda^T(t) \end{bmatrix} = \begin{bmatrix} x_t \\ -\frac{\partial g}{\partial x^T} \Big|_t \end{bmatrix} \quad \triangleright \text{Construct the augmented ODE using PyTorch}$
- 3: **Initialize** neural network parameters θ
- 4: **repeat**
- 5: $V = |\hat{x} - x|^2 \quad \triangleright \text{Forward solve to get the loss } V$
- 6: $\hat{x}_{aug} = \text{ODESolve}(f_{aug}, x_{aug}(0), T, t_0)$
- 7: $\frac{\partial V}{\partial \theta^T} = \lambda^T(0) \frac{\partial x}{\partial \theta^T} \Big|_0 + \int_0^t \lambda^T \frac{\partial f}{\partial \theta^T} dt \quad \triangleright \text{Construct the gradient wrt loss}$
- 8: $\theta \leftarrow \theta - \eta \nabla_{\theta} V$
- 9: **until** convergence criteria is met
- 10: **Return** trained model f_{NN}



Concentration profiles of species A, B, and C in a batch reactor. The NODE model also does quite a good job to predict the concentration profiles. In fact for such simple problem, we hardly see any difference between ResNet and NODE ¹².

¹²All the case studies in this report are built using TorchdiffEq by Chen (2018)

Generative latent model using CNF

- **Latent** = States in state-space model, since we usually cannot measure them directly. Thus, we construct state-estimators.
- **Generative latent** = State-space model that also propagates uncertainty in estimates i.e. you can sample and *generate* multiple trajectories after learning.
- Surprise! Good old Kalman filter can be called a generative latent model! ¹³
- Consider a simple linear transform:

$$x \in \mathcal{N}(m, \sigma^2 I)$$

$$y = Ax$$

$$y \in \mathcal{N}(Am, A\sigma^2 A^T)$$

- Bayesian approach: $p(y|x) \sim p(x|y)p(y)$. Assuming an **uniform prior** the likelihood problem is

$$\theta = \arg \max_{\theta} \log p(x|y; \theta) \quad p(x|y) = A^{-1}y \sim \mathcal{N}(m, \sigma^2 I)$$

- Ofcourse, A is assumed to be non-singular (i.e. unique mapping between x and y ; *bijective mapping*).

¹³Refer Rawlings et al. (2020) for derivation of the Kalman filter using Bayesian estimation principles.

Extending the bayesian framework for any non-linear transform

- Consider following series of transformations (*planar-flow*):

$$x^+ = f_{NN}(x, u; \theta) \quad x(0) \sim \mathcal{N}(m, \sigma^2 I) \quad x(t) \sim ?$$


- Use *change of variables*, compute the **determinant** of Jacobian ($\sim N(o^3)$) and conserve the probability mass as well as maximize the likelihood.
- Instead use continuous framework (*continuous normalizing flow*)¹⁴:

$$\frac{dx}{dt} = f_{NN}(x, u; \theta) \quad x(0) \sim \mathcal{N}(m, \sigma^2 I) \quad x(t) \sim ?$$
$$\frac{\partial p(x)}{\partial t} = -\text{tr} \nabla_x f_{NN}$$

- Cost to get the trace of Jacobian is now linear $\sim N(o)$


¹⁴Refer Chen et al. (2018) for proof

Generative latent model for kinetics



$$\begin{bmatrix} \dot{c}_A \\ \dot{c}_B \\ \dot{c}_C \\ \log p \end{bmatrix} = \begin{bmatrix} -k_1 c_A + k_{-1} c_B \\ 2k_1 c_A + 2k_{-1} c_B - k_2 c_C \\ k_2 c_B \\ -\text{tr} \left(\frac{\partial f_{NN}}{\partial c^T} \right) \end{bmatrix}; \quad \begin{bmatrix} c_A(0) \\ c_B(0) \\ c_C(0) \\ \log p(0) \end{bmatrix} = \begin{bmatrix} 1 \\ 0 \\ 0 \\ \log \mathcal{N}([c_A, c_B, c_C]^T, \sigma^2 I) \end{bmatrix}$$

$$\min_{\theta} V(c) = -\log p(c) + |c - \tilde{c}|^2$$



Robustness analysis of ODEs. We see that the solution c_{B2} is sandwiched between two perturbed solutions c_{B1}, c_{B3} , though all approach the same steady-state. Also note that the slightly time-shifted solution appears as a perturbed solution ¹⁵

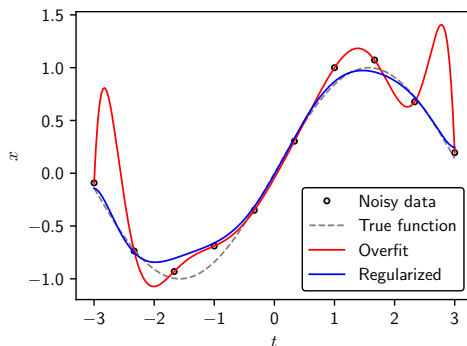
Theoretical bound

$$|x(t+T) - x(t)| \leq \left| \int_t^{t+T} f_{NN}(x, u; \theta) dt \right|$$

¹⁵Refer Ascher and Petzold (1998) for well-posedness theorem of IVPs.

¹⁶Refer Yan et al. (2019)

How to train your NODE?



If we train based only on prediction loss, the model can overfit. However, if we penalize the Jacobian term ($\frac{\partial f_{NN}}{\partial x^T}$), we can regularize the model to damp the oscillations ¹⁷

¹⁷Finlay et al. (2020)

Access the code for all figures at https://github.com/dakeprithvi/2025a_cnf

- U. N. Ascher and L. R. Petzold. *Computer Methods for Ordinary Differential Equations and Differential-Algebraic Equations*. SIAM, Philadelphia, 1998.
- R. T. Chen, Y. Rubanova, J. Bettencourt, and D. K. Duvenaud. Neural ordinary differential equations. *Adv. Neural Inf. Process. Syst.*, 31, 2018.
- R. T. Q. Chen. torchdiffeq, 2018. URL <https://github.com/rtqichen/torchdiffeq>.
- C. Finlay, J.-H. Jacobsen, L. Nurbekyan, and A. Oberman. How to train your neural ode: the world of jacobian and kinetic regularization. In *International conference on machine learning*, pages 3154–3164. PMLR, 2020.
- Y. Liu, A. Sholokhov, H. Mansour, and S. Nabi. Physics-informed koopman network for time-series prediction of dynamical systems. In *ICLR 2024 Workshop on AI4DifferentialEquations In Science*, 2024.
- L. Ljung. *System Identification: Theory for the User*. Prentice Hall, New Jersey, second edition, 1999.
- M. Raissi, P. Perdikaris, and G. E. Karniadakis. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *J. Comput. Phys.*, 378:686–707, 2019.

- J. B. Rawlings and J. G. Ekerdt. *Chemical Reactor Analysis and Design Fundamentals*. Nob Hill Publishing, Santa Barbara, CA, 2nd, paperback edition, 2020. 664 pages, ISBN 978-0-9759377-4-7.
- J. B. Rawlings, D. Q. Mayne, and M. M. Diehl. *Model Predictive Control: Theory, Computation, and Design*. Nob Hill Publishing, Santa Barbara, CA, 2nd, paperback edition, 2020. 770 pages, ISBN 978-0-9759377-5-4.
- B. Sengupta, K. J. Friston, and W. D. Penny. Efficient gradient computation for dynamical models. *NeuroImage*, 98:521–527, 2014.
- A. Sholokhov, Y. Liu, H. Mansour, and S. Nabi. Physics-informed neural ode (pinode): embedding physics into models using collocation points. *Scientific Reports*, 13(1): 10166, 2023.
- J. V. Villadsen and W. E. Stewart. Solution of boundary-value problems by orthogonal collocation. *Chem. Eng. Sci.*, 22:1483–1501, 1967.
- S. Wang, S. Sankaran, and P. Perdikaris. Respecting causality is all you need for training physics-informed neural networks, 2022. URL <https://arxiv.org/abs/2203.07404>.
- H. Yan, J. Du, V. Y. Tan, and J. Feng. On robustness of neural ordinary differential equations. *arXiv preprint arXiv:1910.05513*, 2019.
- J. Zhuang, N. Dvornik, X. Li, S. Tatikonda, X. Papademetris, and J. Duncan. Adaptive checkpoint adjoint method for gradient estimation in neural ode. In *International Conference on Machine Learning*, pages 11639–11649. PMLR, 2020.